




**Figure 1.** Mean classification error vs. query budget  $m$  for three selection strategies across three tasks ( $n=80$ ). Stepwise selection matches uniform sampling accuracy with 5–10 $\times$  fewer queries.

## 1 Motivation

Classifying black-box generative models requires querying each model and comparing responses via pairwise energy distances. In a typical workflow, a labeled pilot set of  $n$  models is queried on a large pool of  $M$  queries, and the cached responses are used to estimate the discriminative structure. The cost of classifying a *new* model scales with the number of queries  $m$  it must answer — each query is an API call. Identifying a small, high-quality query set  $Q^*$  of size  $m \ll M$  that preserves classification accuracy therefore directly reduces the per-model audit cost for all future models.

## 2 Methods

**1. Uniform Sampling** (baseline). Sample  $m$  queries uniformly at random from the pool of  $M$  queries. No pilot information is used.

**2. Estimated Signal Sampling.** Fit a two-component GMM to the SVD loadings  $|U_{q,\ell}|$  of the energy tensor  $E$ . Queries assigned to the higher-mean component form the estimated signal set  $\hat{S}$ . Sample  $m$  queries uniformly from  $\hat{S}$ . This exploits zero-set structure but ignores query redundancy.

**3. Forward Stepwise Selection.** Define a response  $z_{ij} = \mathbf{1}[y_i \neq y_j]$  for each model pair and a predictor  $E_{q,ij} = \|g(f_i(q)) - g(f_j(q))\|^2$  for each query  $q$ . At each step, add the query whose  $E_{q,\cdot}$  has the largest partial correlation with the residual of  $z$  after projecting out previously selected queries. This is Gram-Schmidt orthogonalization on  $E$  w.r.t.  $z$ , ensuring each query contributes maximally non-redundant signal. Cost:  $O(m \cdot M \cdot \binom{n}{2})$  — negligible.

## 3 Setup

| Task        | $n$ | $M$ | Description              |
|-------------|-----|-----|--------------------------|
| Motivating  | 100 | 200 | Sensitive training data  |
| Sys. Prompt | 100 | 200 | Covert persuasion bias   |
| RAG         | 120 | 200 | Unauthorized doc. stores |

$n_{\text{train}} = 80$  models for estimation and classification; rest held out. MDS (8 components) + LDA. Averaged over 100 random splits.

## 4 Results

Stepwise dominates at small  $m$ : at  $m=2$ , stepwise error is 10% (motivating) and 11% (system prompt) vs. 28% for uniform. Estimated signal is intermediate, confirming the GMM partition captures real structure. On the system prompt task, stepwise at  $m=5$  matches uniform at  $m=50$  — a 10 $\times$  reduction in per-model query cost.